Retained linear-output RMS | fixed first 256 TRAIN features | upper row: initialization; lower row: after one epoch

PRODUCER-RETAINED forward summaries; independently checked for provenance/finite metadata, NOT independently recomputed feature inference.

Input/output maps are Dense with bias; all internal maps are bias-free. Final output RMS includes output bias and does not establish useful discrimination.
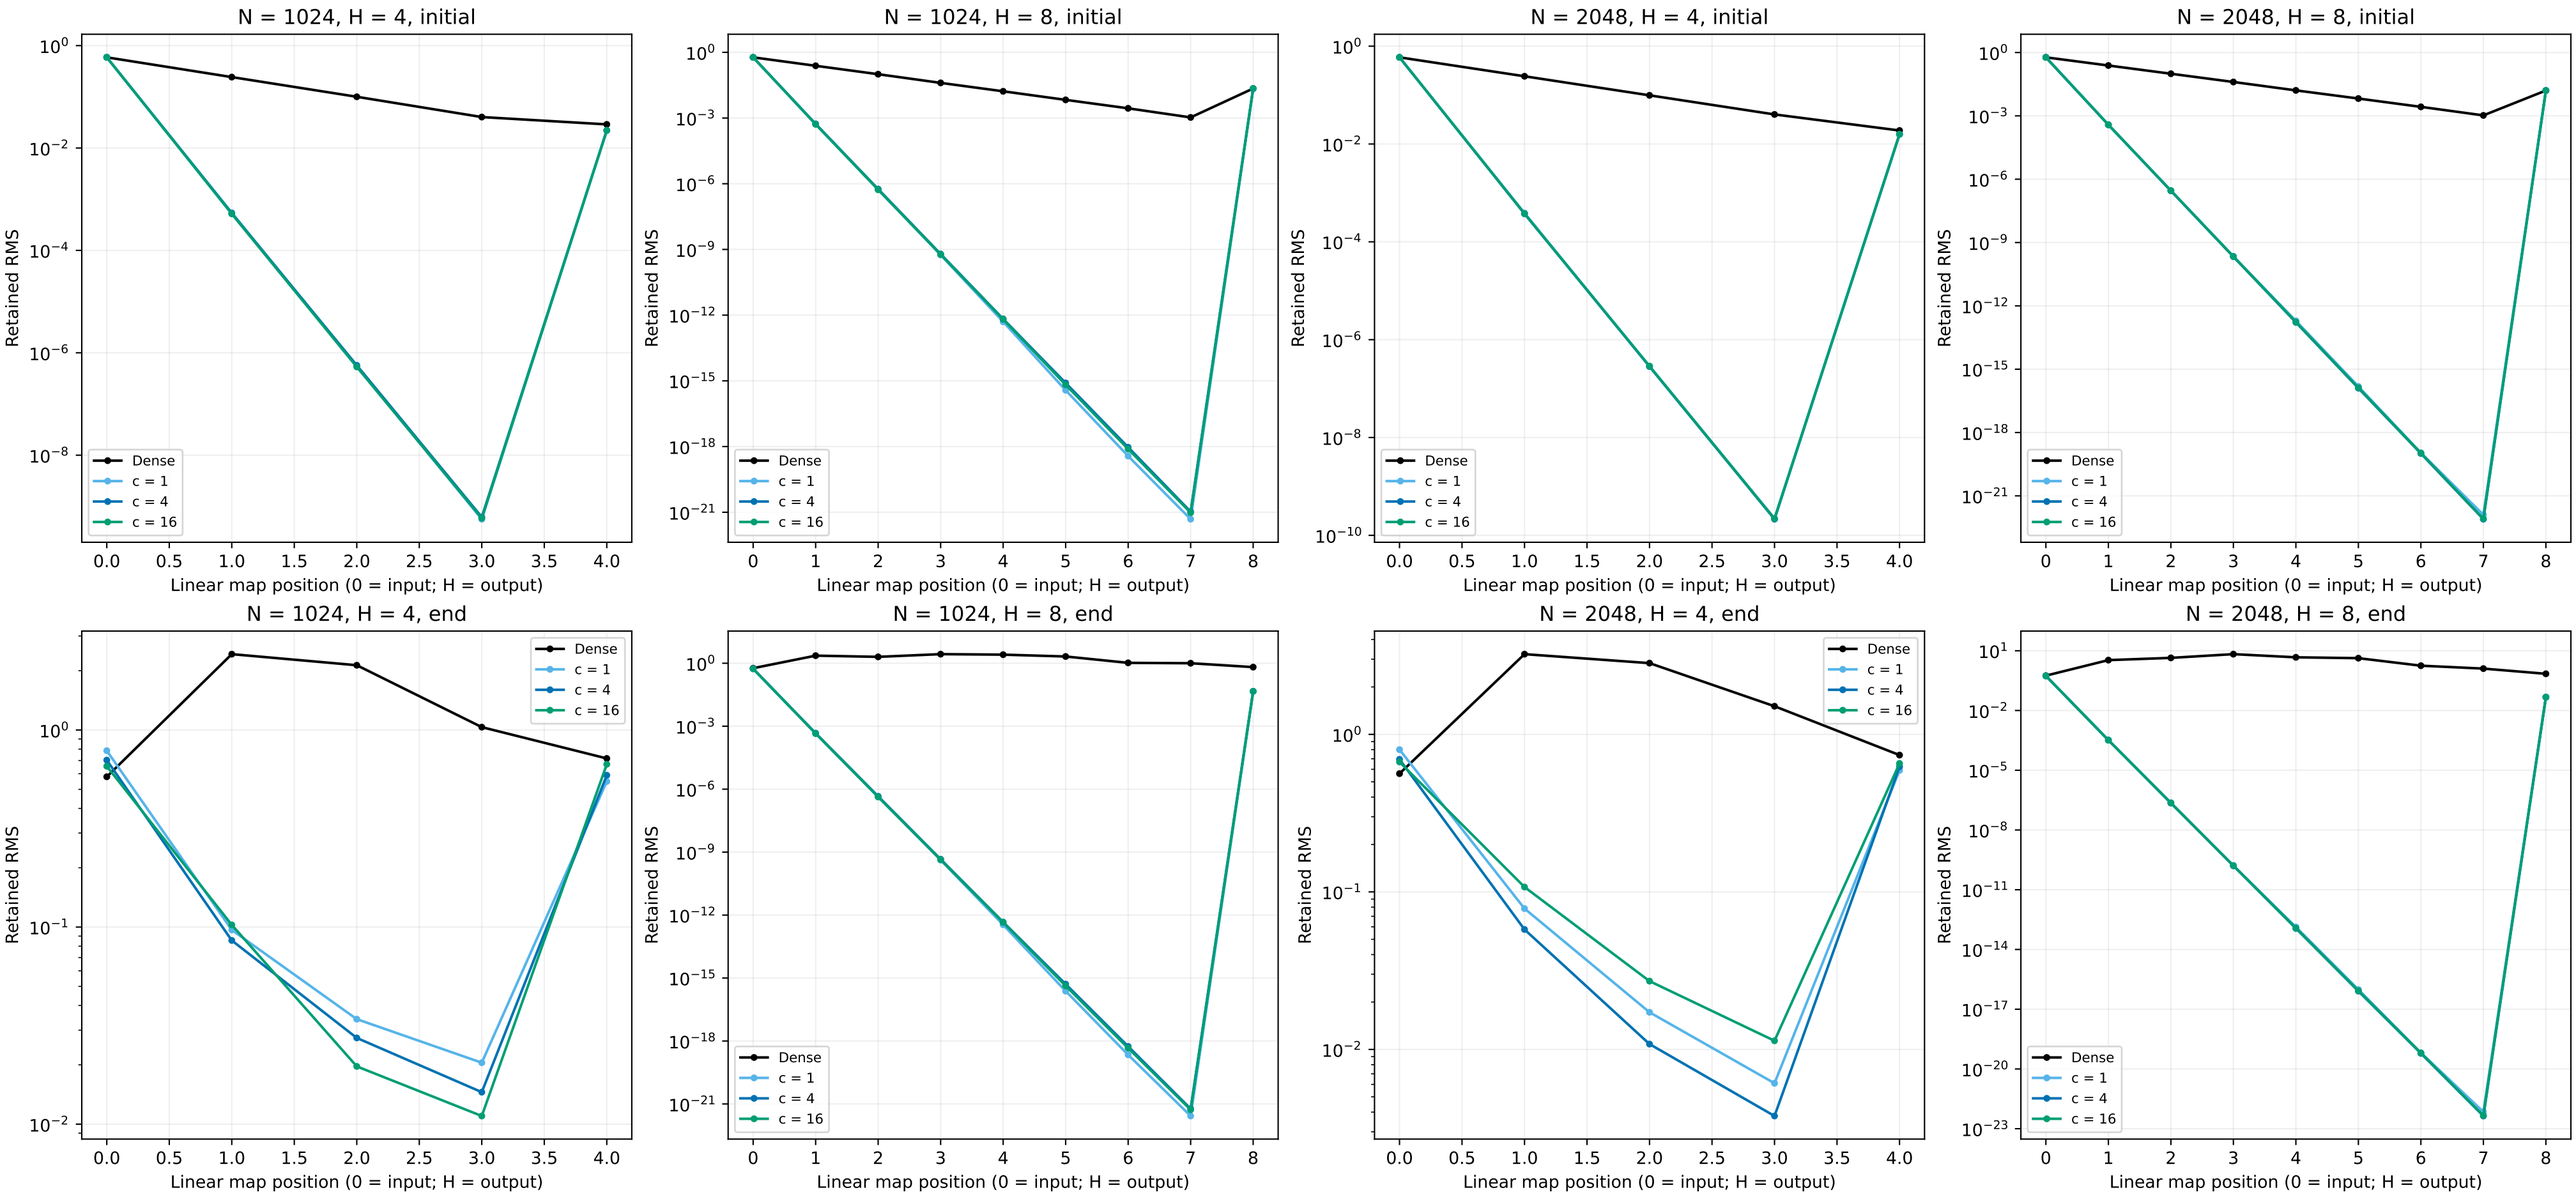
Log scale: exact zeros are masked without an epsilon floor. These are depth profiles, not learning/convergence trajectories or proof of a unique failure cause.