Stock NONE / unit gains / legacy initialization | seed 17 | one full 100,000-row epoch | 25,000 B4 updates per case | eager LR 0.0003

Selected TRAIN and DEV AUC independently recomputed from saved predictions. Endpoints only: no fabricated convergence curve, HPO, formal NI or minimum-Atom claim.

All 16 cases ran in one RTX 4090 cohort. Training clock excludes preparation/evaluation/diagnostics/checkpoint work; allocated/reserved peaks were captured before evaluation.
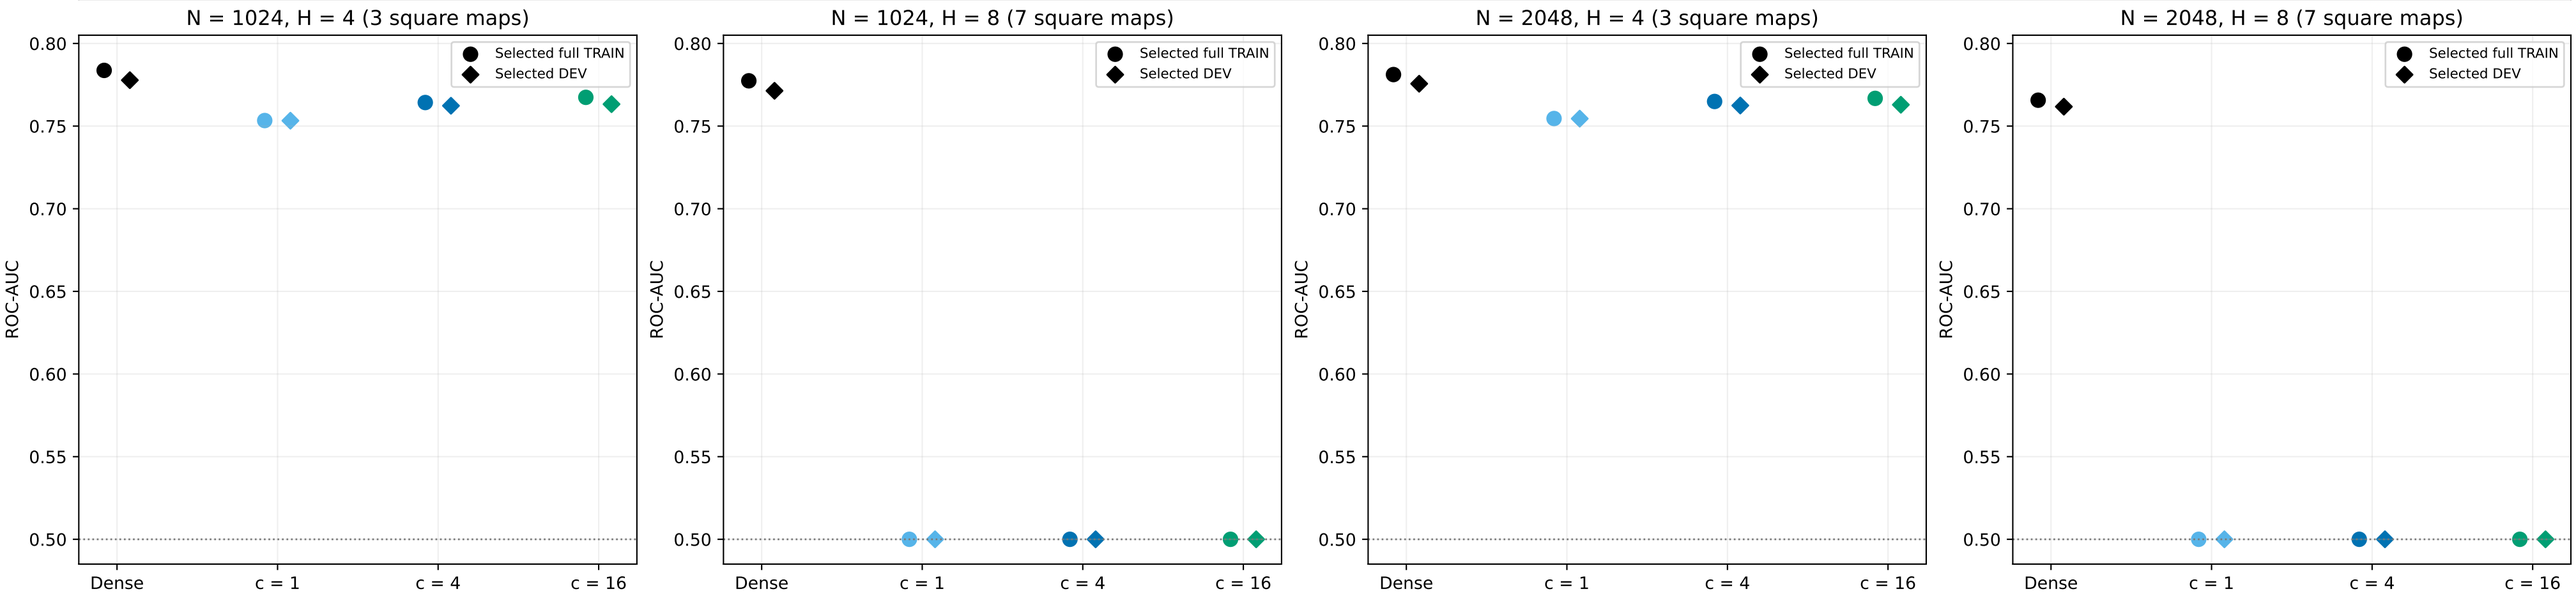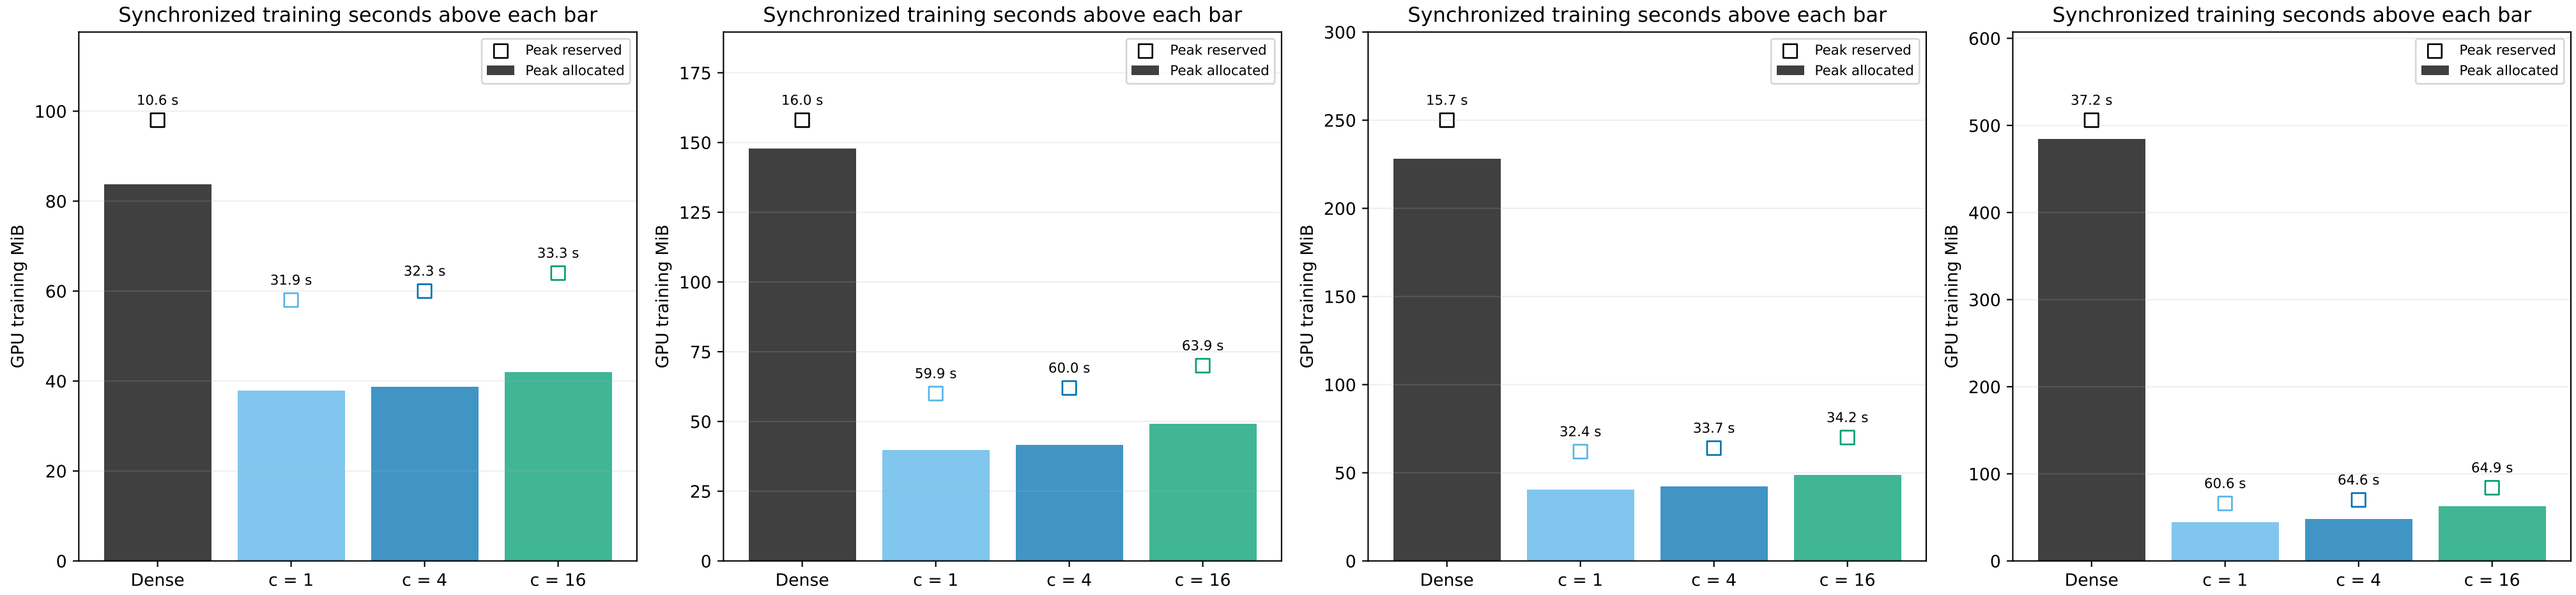
GPU training peaks include resident TRAIN/DEV, model/optimizer/native buffers; reserved cache can persist. CPU-reader RSS is separate. No matched-quality or normalization-necessity claim.